

# Multi-Core Programming

Revision

# Multi-core Programming

- Parallel Programming Models
  - Message Passing
  - Shared Address Space
  - Data Parallel
- Execution Models
  - Cluster
  - Symmetric Multiple Processors (or Multi-core)
  - Single Instruction Stream in Single core
- Performance

# Performance Models

- Instruction execution time (cycles, response time)
- Throughput when multi-core

# Improving Performance

- Hardware Tech.
  - Cache
  - Hardware multi-threading
  - Prefetching
  - Pipelining
  - ...
- Programming

# Programming for Performance

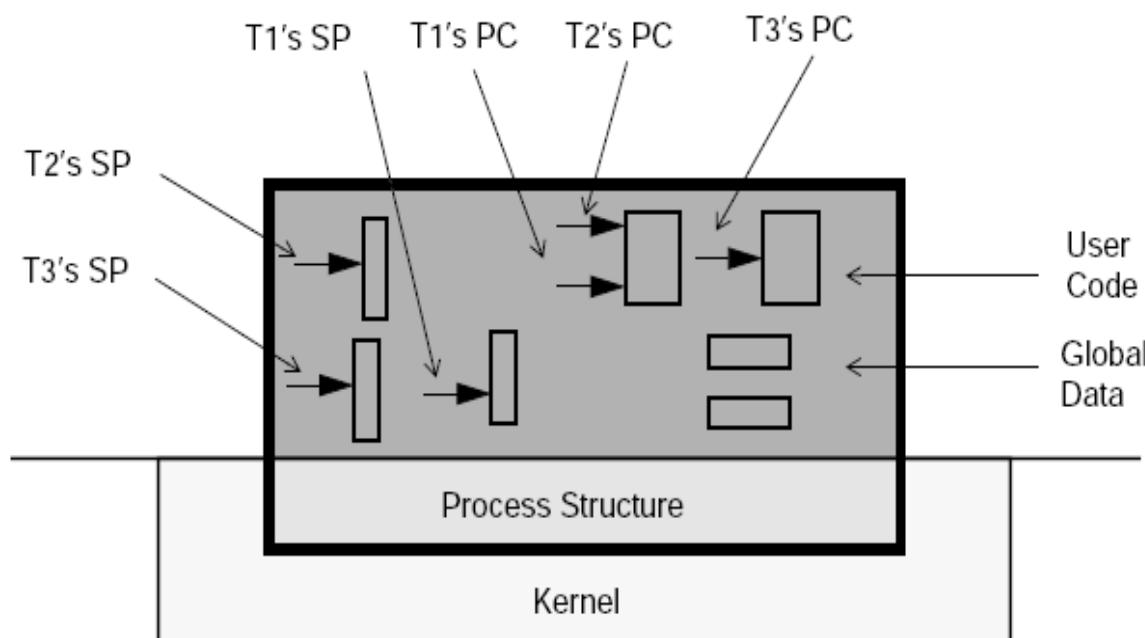
- Good Algorithms
- Utilization of Runtime and System
- Utilization of Computation Units
- Workload Balance
- Memory access
- Communication
- Synchronization
- Scheduling and Mapping
- ...

# Lecture02.Multithreading

- Thread computational model layers:
  - User level threads
  - Kernel level threads
  - Hardware threads

The diagram illustrates the thread computational model layers. It shows three levels: 'User level threads' (defining threads), 'Kernel level threads' (operating threads), and 'Hardware threads' (executing threads). Red arrows point from left to right between these levels. To the right, a vertical bracket groups all three levels under the label 'flow of threads'.
- Mapping
  - Thread Level Scheduler in user space
  - OS scheduler in Kernel space

# Explain what is thread according to the following diagram.



- Multiple threads in a process
- All threads share the global data of the process
- Each thread has its own program counter and stack point
- All stacks of threads are all within the memory space of the process

# Lecture03.Parallel Programming Models

- Programming Abstraction and Implementation
- 3 Programming Models
  - Shared Address Space Model
  - Message Passing Model
  - Data Parallel Model
- 3 Architectures
  - Multi-core
  - Cluster
  - SIMD

Why Gather and scatter operations sometimes are needed for data parallel communication? How can these operations be implemented in an old ISPC code?

- Because SIMD Units require aligned contiguous data
- ISPC may use indirect reference to gather or scatter
- Gather is supported in AVX2 and scatter exists in AVX512

# Lecture04.Message Passing Model

- Non-Buffered Blocking Message Passing Operations
- Buffered Blocking Message Passing Operations
- Non-Blocking Message Passing Operations
- Process Grouping: Communicators
- MPI Point to Point Communication
- MPI Collective Communication
  
- Simple programs

What things you have to ensure when you write a program with non-blocking message passing operations ?

- The **programmer** must ensure semantics of the send and receive.

MPI uses Communicator to group processes. Why commuincator is important in MPI?

- A **communicator** defines a *communication domain* - a set of processes that are allowed to communicate with each other.
- MPI Collective communications are all based on the communicator they belong.

# Lecture05.SAS-PTthreads

- POSIX Thread Creation and Exit
- *Thread Attributes*
- *Thread Synchronization*
- Write simple parallel programs using synchronizations

# **Lecture07.OpenMP - Components of OpenMP**

- **The OpenMP Execution Model**
- **Components of OpenMP**
  - Directives
  - Environment Variables
  - Runtime
- **Parallel Region**
- **Work-sharing constructs**
- **Tasks**
- **Synchronization and Scheduling**

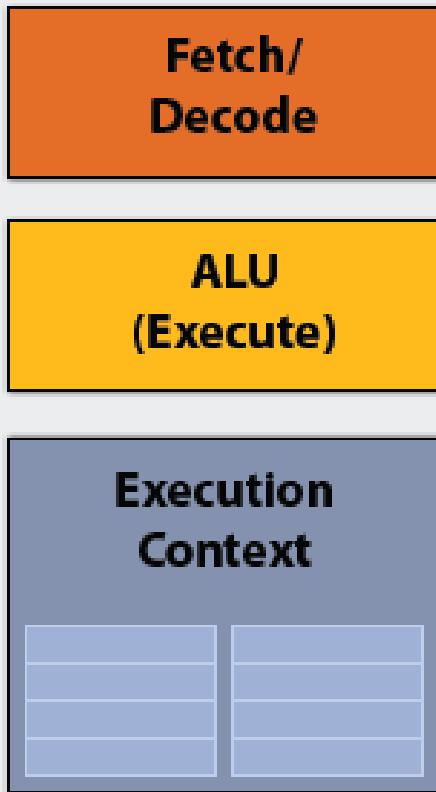
# Lecture09.ISPC Parallel Programming Model

- Gang of ISPC program instances
- SIMD Implementation
  - Assignment
- ISPC Tasks

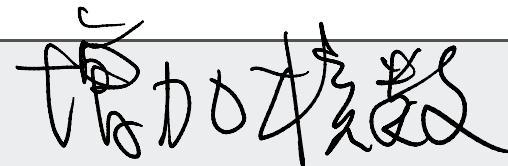
# Lecture11.Mixed Parallel Programming

- MPI+OpenMP Hybrid Model
- Usually MPI+X

# Lecture12.Modern Processor



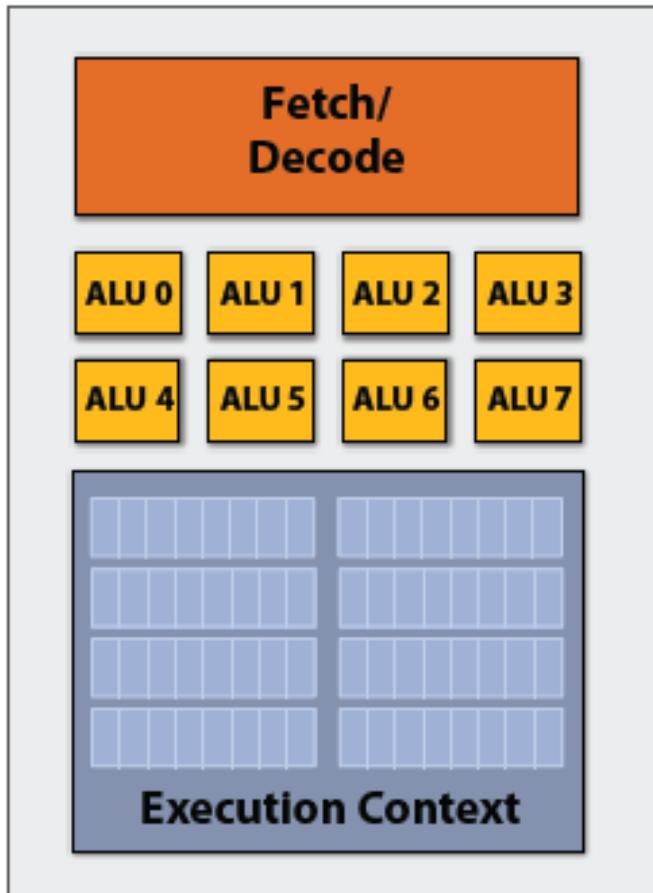
Idea #1:



**Use increasing transistor count to add more cores to the processor**

**Rather than use transistors to increase sophistication of processor logic that accelerates a single instruction stream (e.g., out-of-order and speculative operations)**

# Add ALUs to increase compute capability



Idea #2:

Amortize cost/complexity of managing an instruction stream across many ALUs

## SIMD processing

Single instruction, multiple data

Same instruction broadcast to all ALUs  
Executed in parallel on all ALUs

向量程序

支持向量的内核函数

# Vector program (using AVX intrinsics)

```
#include <iimmintrin.h>

void sinx(int N, int terms, float* x, float* result)
{
    float three_fact = 6; // 3!
    for (int i=0; i<N; i+=8)
    {
        __m256 origx = _mm256_load_ps(&x[i]);
        __m256 value = origx;
        __m256 numer = _mm256_mul_ps(origx, _mm256_mul_ps(origx, origx));
        __m256 denom = _mm256_broadcast_ss(&three_fact);
        int sign = -1;

        for (int j=1; j<=terms; j++)
        {
            // value += sign * numer / denom
            __m256 tmp = _mm256_div_ps(_mm256_mul_ps(_mm256_set1ps(sign), numer), denom);
            value = _mm256_add_ps(value, tmp);

            numer = _mm256_mul_ps(numer, _mm256_mul_ps(origx, origx));
            denom = _mm256_mul_ps(denom, _mm256_broadcast_ss((2*j+2) * (2*j+3)));
            sign *= -1;
        }
        _mm256_store_ps(&result[i], value);
    }
}
```

Intrinsics available to C programmers

跨度8个float

single precise float

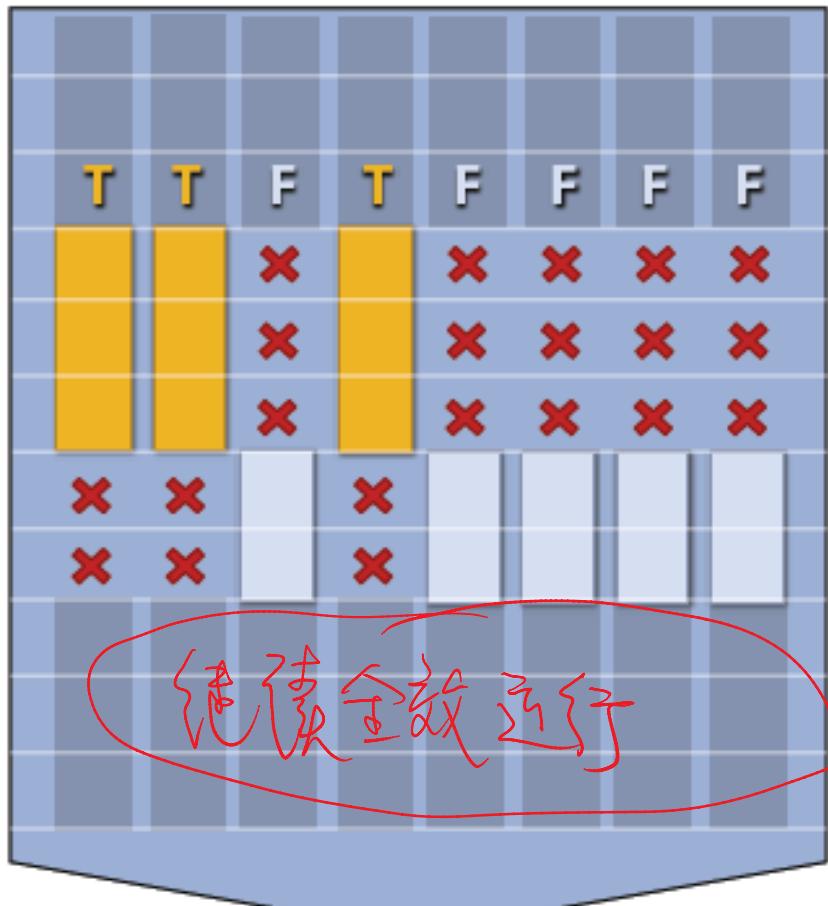
$32 \times 8 = 256$  跨度8个float

同时8个乘(向量乘)

广播3!

# After branch: continue at full performance

Time (clocks)



(assume logic below is to be executed for each element in input array 'A', producing output into the array 'result')

<unconditional code>

```
float x = A[i];  
if (x > 0) {  
    float tmp = exp(x,5.f);  
    tmp *= kMyConst1;  
    x = tmp + kMyConst2;  
} else {  
    float tmp = kMyConst1;  
    x = 2.f * tmp;  
}
```

<resume unconditional code>

```
result[i] = x;
```

# Lecture13.Accessing Memory

- **Memory latency**

存儲延遲

- The amount of time for a memory request (e.g., load, store) from a processor to be serviced by the memory system
- Example: 100 cycles, 100 nsec

- **Memory bandwidth**

存儲帶寬

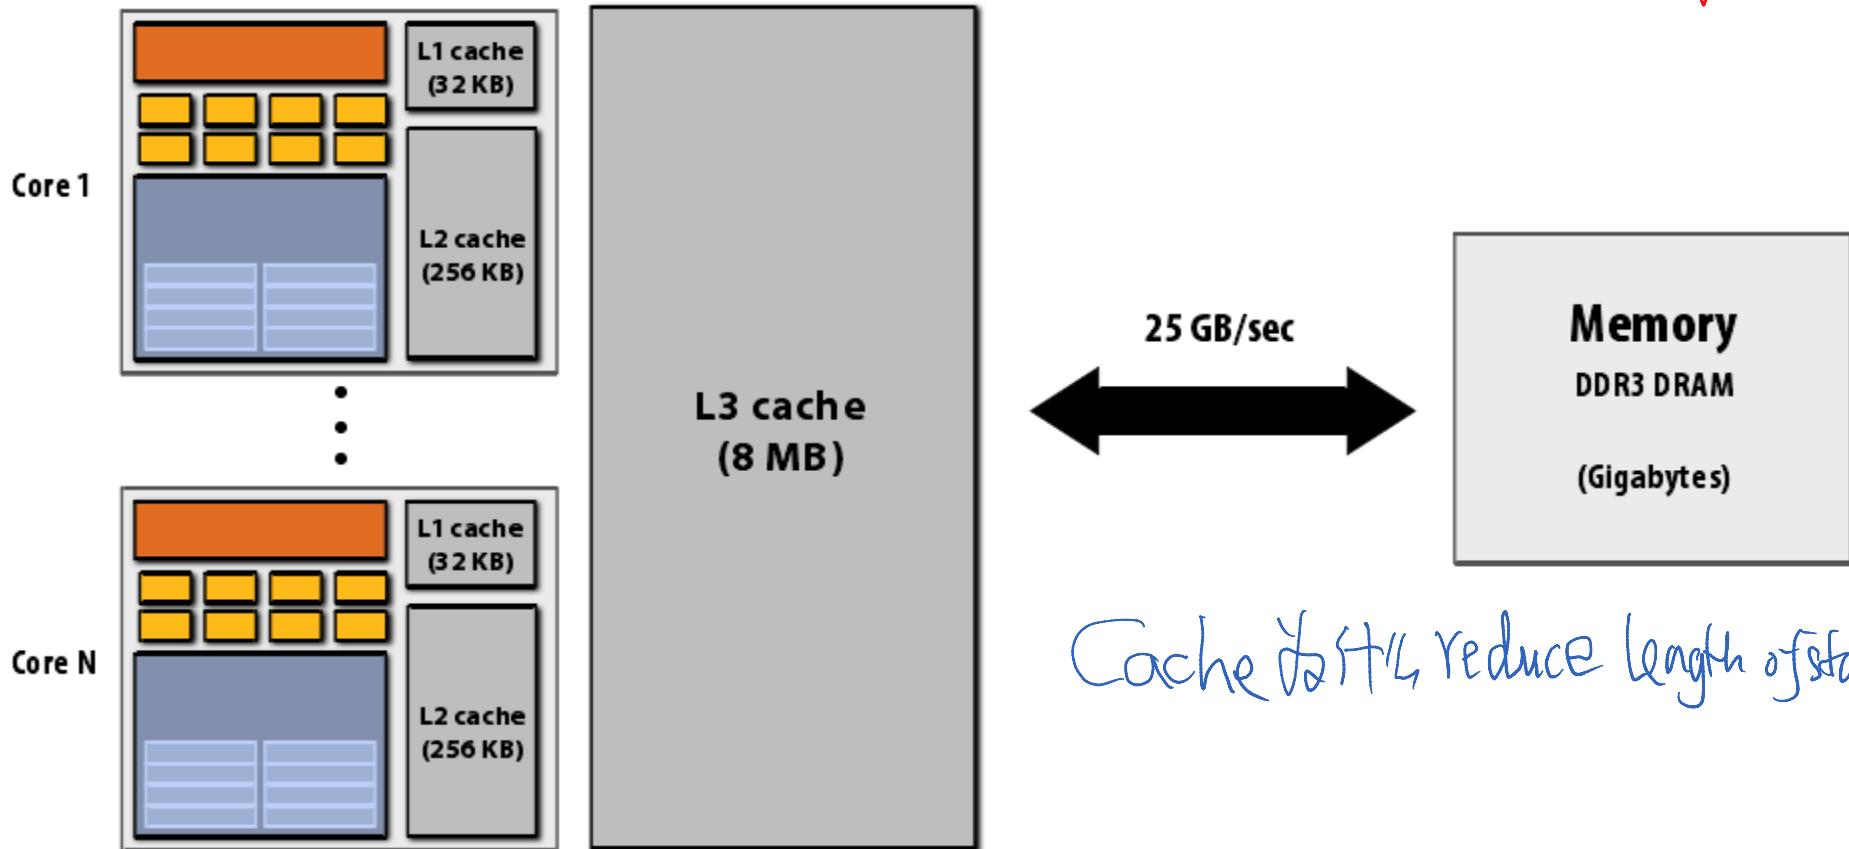
- The rate at which the memory system can provide data to a processor
- Example: 20 GB/s

# Caches reduce length of stalls (reduce latency)

Processors run efficiently when data is resident in caches

Caches reduce memory access latency (\* Caches also provide high bandwidth data transfer to CPU)

缓存降低内存访问延时，提高带宽



# Prefetching reduces stalls (hides latency)

- All modern CPUs have logic for prefetching data into caches
    - Dynamically analyze program's access patterns, predict what it will access soon

硬件处理器都有了取数缓存 [Cache in CPU]

- Reduces stalls since data is resident in cache when accessed

The diagram illustrates a sequence of memory access requests from registers r2 and r3. The requests are labeled as "predict value of r2, initiate load" and "predict value of r3, initiate load". Ellipses indicate more requests. Red arrows point from the text "data arrives in cache" to the initiation points of the requests, indicating that both requests are being serviced by the same cache unit simultaneously.

**Note: Prefetching can also reduce performance if the guess is wrong (hogs bandwidth, pollutes caches)**

**(more detail later in course)**

ld r0 mem[r2]      ]  
ld r1 mem[r3]      ] These loads are cache hits  
add r0, r0, r1

避免了延迟(隐性延迟)(已转Cache中)

(交错) 多线程减少拖延期

## Multi-threading reduces stalls

交错(交替)线程

- Idea: **interleave** processing of multiple threads on the same core to hide stalls
- Like prefetching, multi-threading is a latency **hiding**, not a latency **reducing** technique

这里先搞 interleave 多线程, 后面会讲 SMT。

同核上的交错线程可以掩盖拖延 (隐藏拖延)  
(hide stalls)

# Bandwidth is a critical resource

Performant parallel programs will:

- Organize computation to fetch data from memory less often  
*不要那么频繁地取数据。*
  - Reuse data previously loaded by the same thread (traditional intra-thread temporal locality optimizations)
  - Share data across threads (inter-thread cooperation)
- Request data less often (instead, do more arithmetic: it's "free")  
*不要那么快(频繁)地申请数据*
  - Useful term: "arithmetic intensity" — ratio of math operations to data access operations in an instruction stream
  - Main point: programs must have high arithmetic intensity to utilize modern processors efficiently

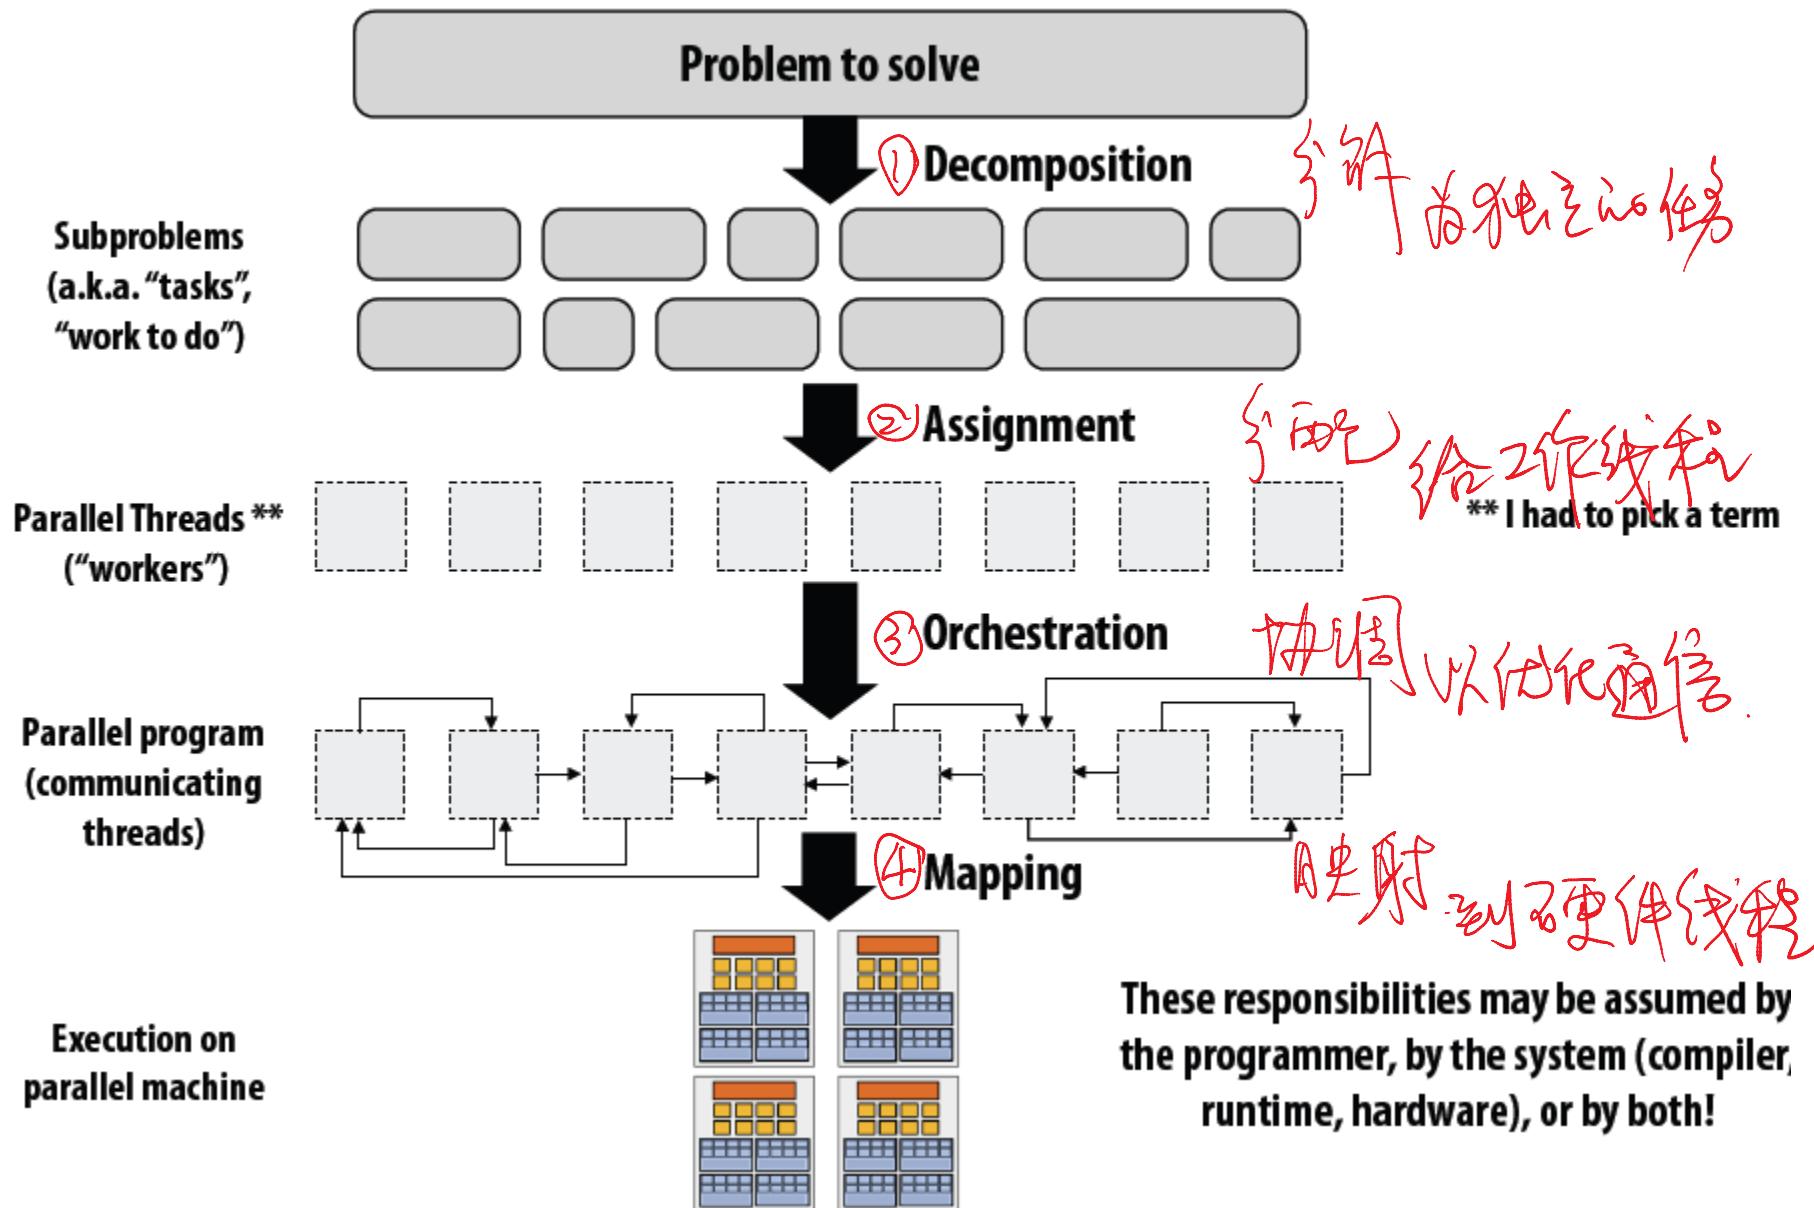
*尽量提高算术强度以充分利用现代处理器  
算术强度与"计算与通信比"类似*

# Lecture14.Programming Basics

- Decomposition
- Assignment
- Orchestration
- Mapping

# 创建并行程序的步骤

# Creating a parallel program



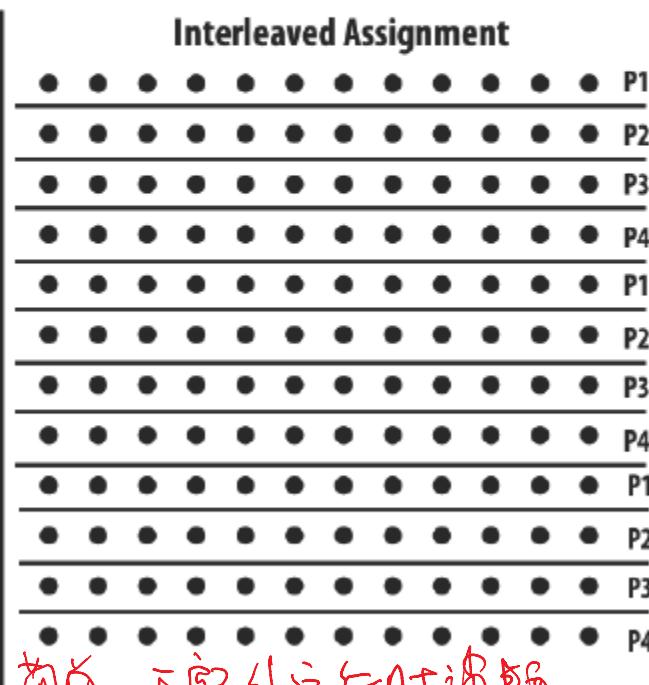
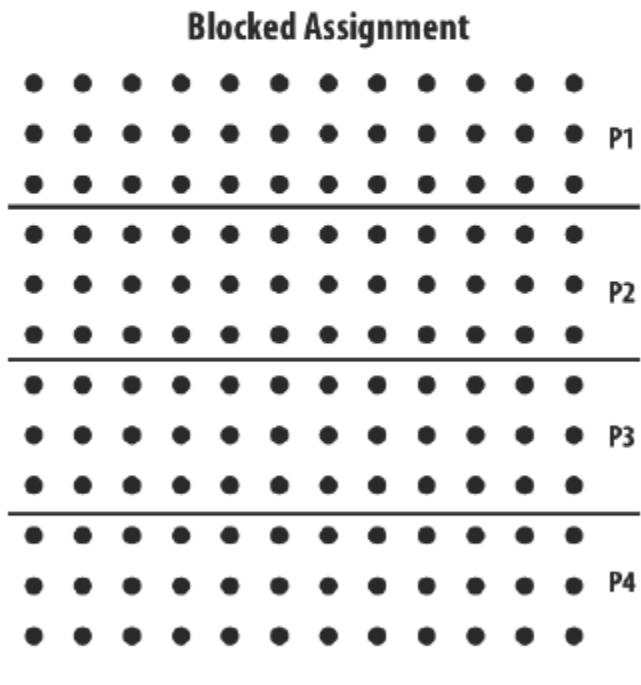
# Lecture15.Workload Balancing

- Task Assignment and Scheduling

# Static assignment

静态分配时对负载  
均衡很考虑。

- Assignment of work to threads is pre-determined 预先分配(但可能在运行时, 依然能动态)
  - Not necessarily determined at compile-time (assignment algorithm may depend on runtime parameters such as input data size, number of threads, etc.)
- Recall solver example: assign equal number of grid cells (work) to each thread (worker)
  - We discussed two static assignments of work to workers (blocked and interleaved)



均匀分配  
线程间同意  
着工作负担

简单, 无额外运行时消耗

Good properties of static assignment: simple, essentially zero runtime overhead  
(in this example: extra work to implement assignment is a little bit of indexing math)

# Dynamic assignment

动态分配

Program determines assignment dynamically at runtime to ensure a well distributed load. (The execution time of tasks, or the total number of tasks, is unpredictable.)

Sequential program  
(independent loop iterations)

```
int N = 1024;
int* x = new int[N];
bool* prime = new bool[N];

// initialize elements of x here

for (int i=0; i<N; i++)
{
    // unknown execution time
    is_prime[i] = test_primality(x[i]);
}
```

执行时间不确定  
因为对x[i]大小的预测，导致执行语句的执行时间相差太大。

Parallel program  
(SPMD execution by multiple threads,  
shared address space model)

```
int N = 1024;
// assume allocations are only executed by 1 thread
int* x = new int[N];
bool* is_prime = new bool[N];

// initialize elements of x here

LOCK counter_lock;
int counter = 0; // shared variable

while (1) {
    int i;
    lock(counter_lock);
    i = counter++;
    unlock(counter_lock);
    if (i >= N)
        break;
    is_prime[i] = test_primality(x[i]);
}
```

①如何分配的？

②为什么要lock？

③为什么动态分配比静态分配好？

④这种分配方法的优点是什么？  
(不锁行不行)。

atomic\_incr(counter);

# Dynamic assignment using a work queue

用工作队列进行动态分配.

Sub-problems  
(a.k.a. "tasks", "work")



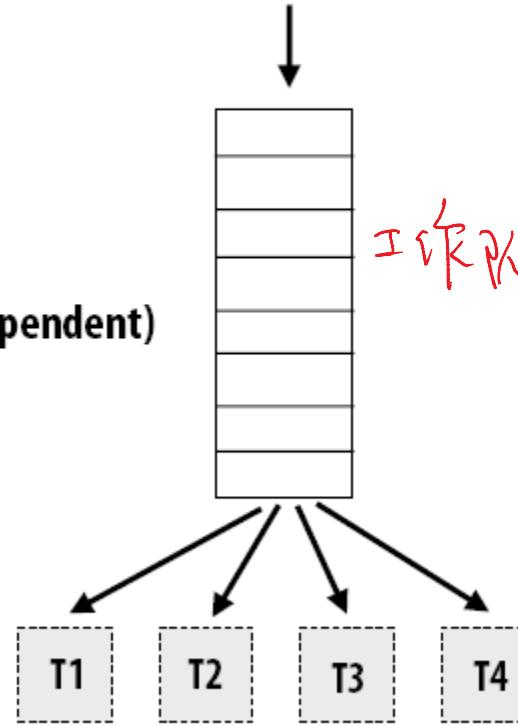
Shared work queue: a list of work to do  
(for now, let's assume each piece of work is independent)

工作队列

Worker threads:

Pull data from shared work queue

Push new work to queue as it is created



执行线程从队列  
取任务也进行执行  
并且可以将新任务加入

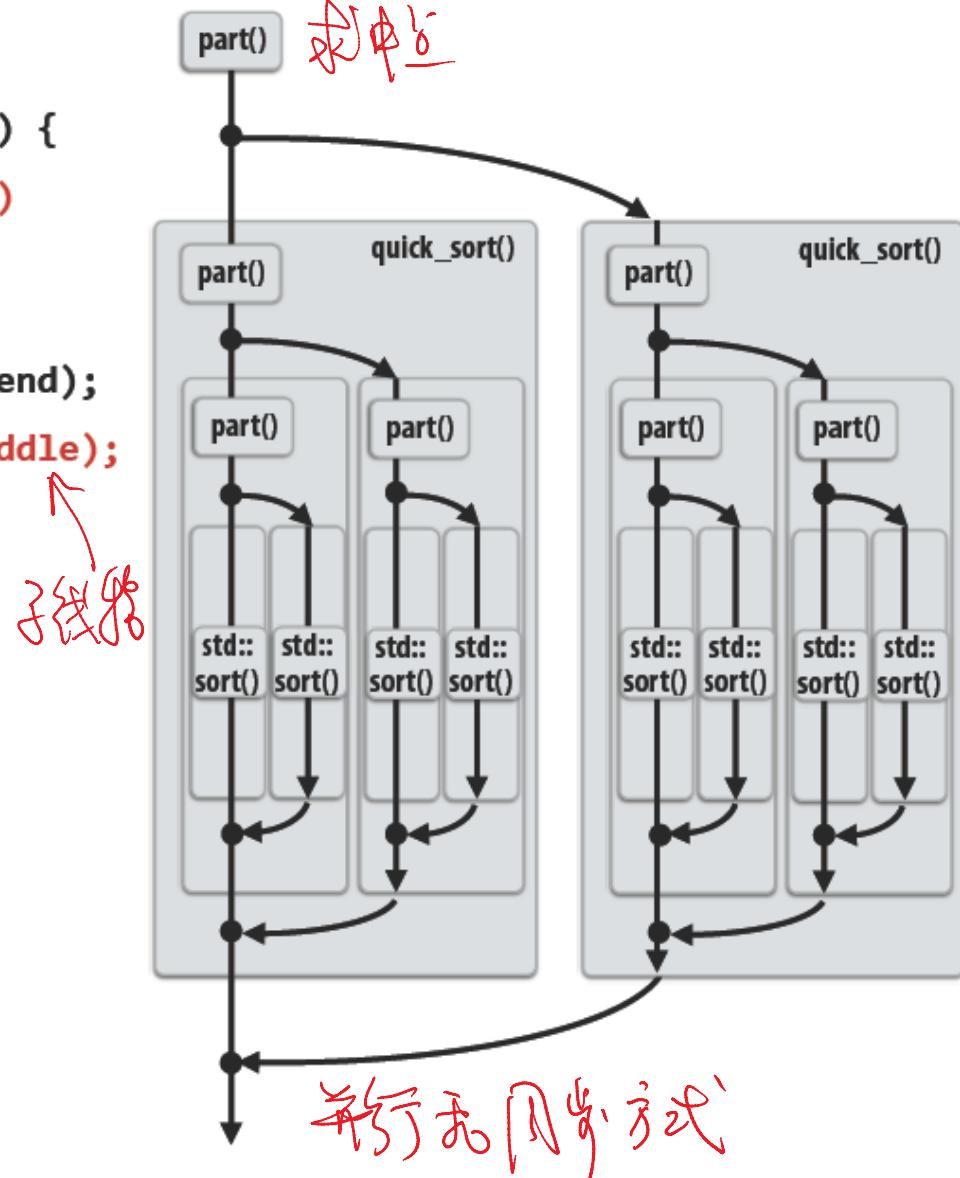
# Lecture16.Scheduling fork-join parallelism

# Parallel quicksort in Cilk Plus (Fork-Join Abstraction)

```
void quick_sort(int* begin, int* end) {  
    if (begin >= end - PARALLEL_CUTOFF)  
        std::sort(begin, end);  
    else {  
        int* middle = partition(begin, end);  
        cilk_spawn quick_sort(begin, middle);  
        quick_sort(middle+1, last);  
    }  
}
```

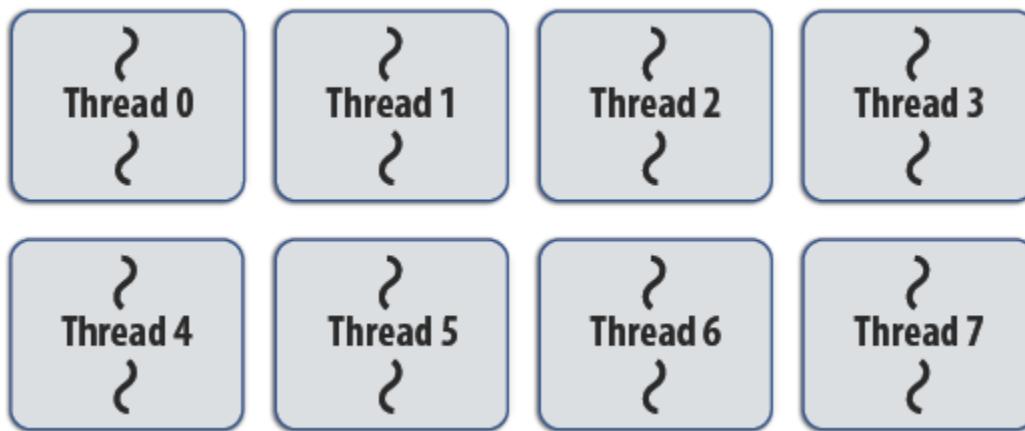
Sort sequentially if problem size is sufficiently small (overhead of spawn trumps benefits of potential parallelization)

当Size太小时，用串行sort.



# 运行时了{之上实现 Pool of worker threads 2个线程池

- The Cilk Plus runtime maintains pool of worker threads
  - Think: all threads created at application launch \*  
\* It's perfectly fine to think about it this way, but in reality, runtimes tend to be lazy and initialize worker threads on the first Cilk spawn. (This is a common implementation strategy, ISPC does the same with worker threads that run ISPC tasks.)
  - Exactly as many worker threads as execution contexts in the machine 2个线程池与机器的执行环境数相同



Example: Eight thread worker pool for my quad-core laptop with Hyper-Threading

四核+HT每物理可用8个线程池

```
while (work_exists()) {  
    work = get_new_work();  
    work.run();  
}
```

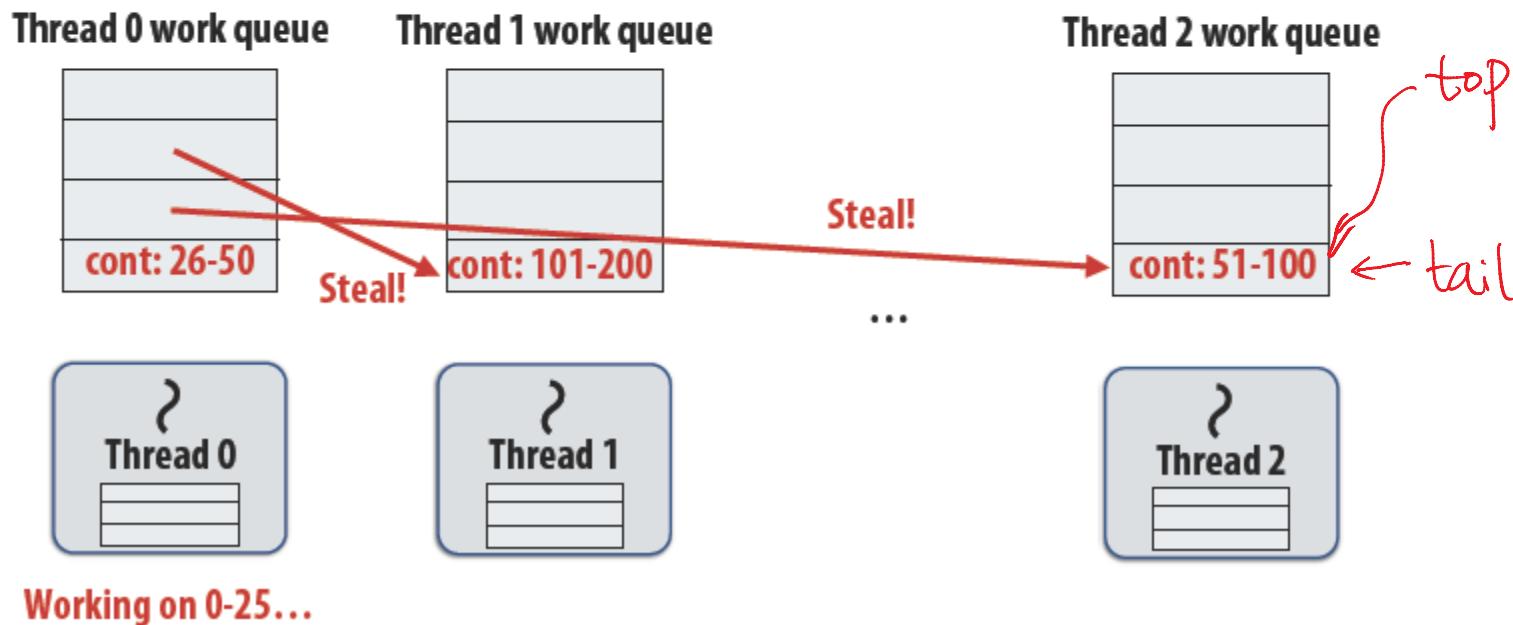
运行2线程池

# Implementing work stealing: dequeue per worker

用双向队列实现盗贼。

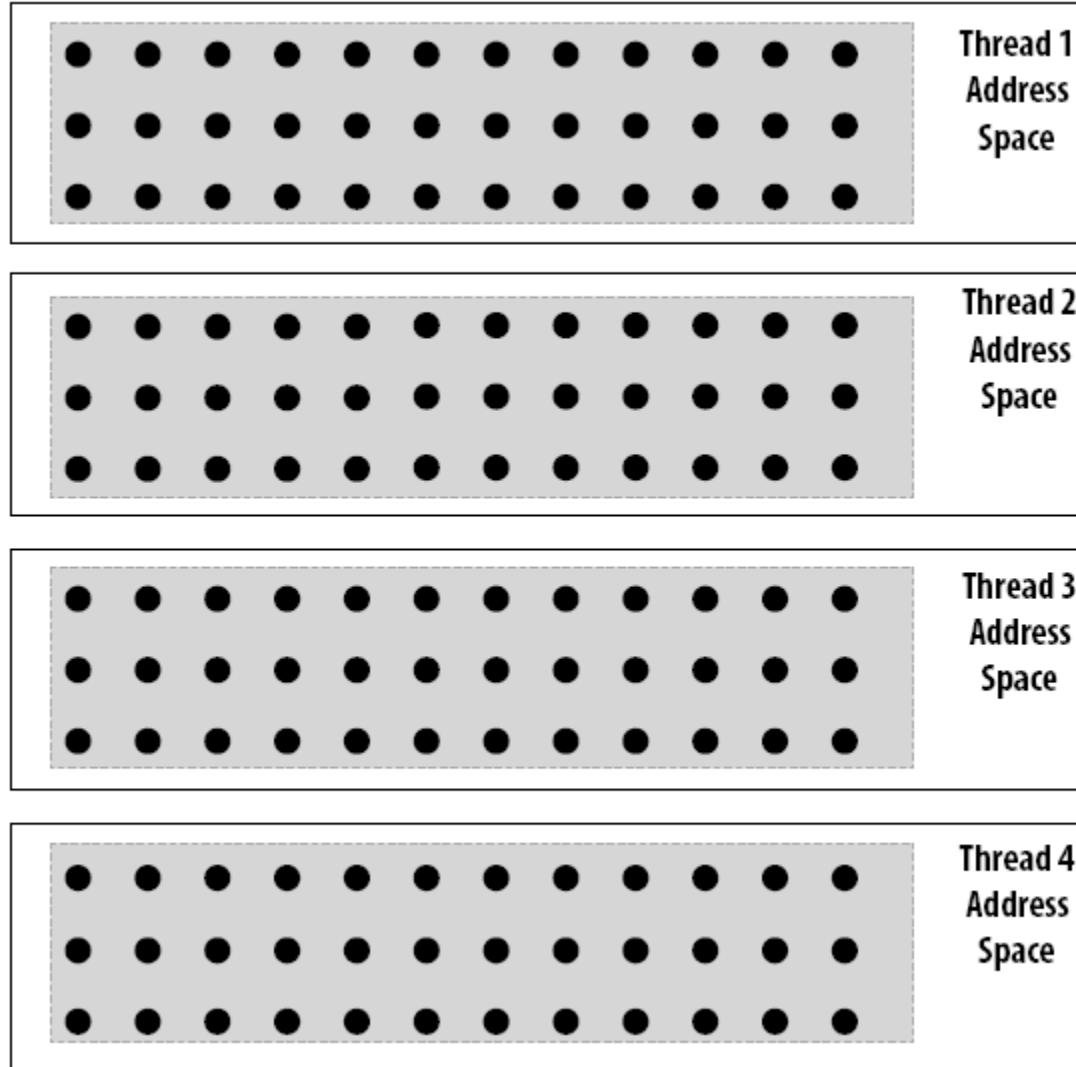
## Work queue implemented as a deque (double ended queue)

- Local thread pushes/pops from the “tail” (bottom) 本地线程从队尾操作(取,入)
- Remote threads steal from “head” (top) 其它线程从队头操作(取)  
为什么要这样操作?



# **Lecture17.Communication Cost**

# Message passing model: each thread operates in its own address space

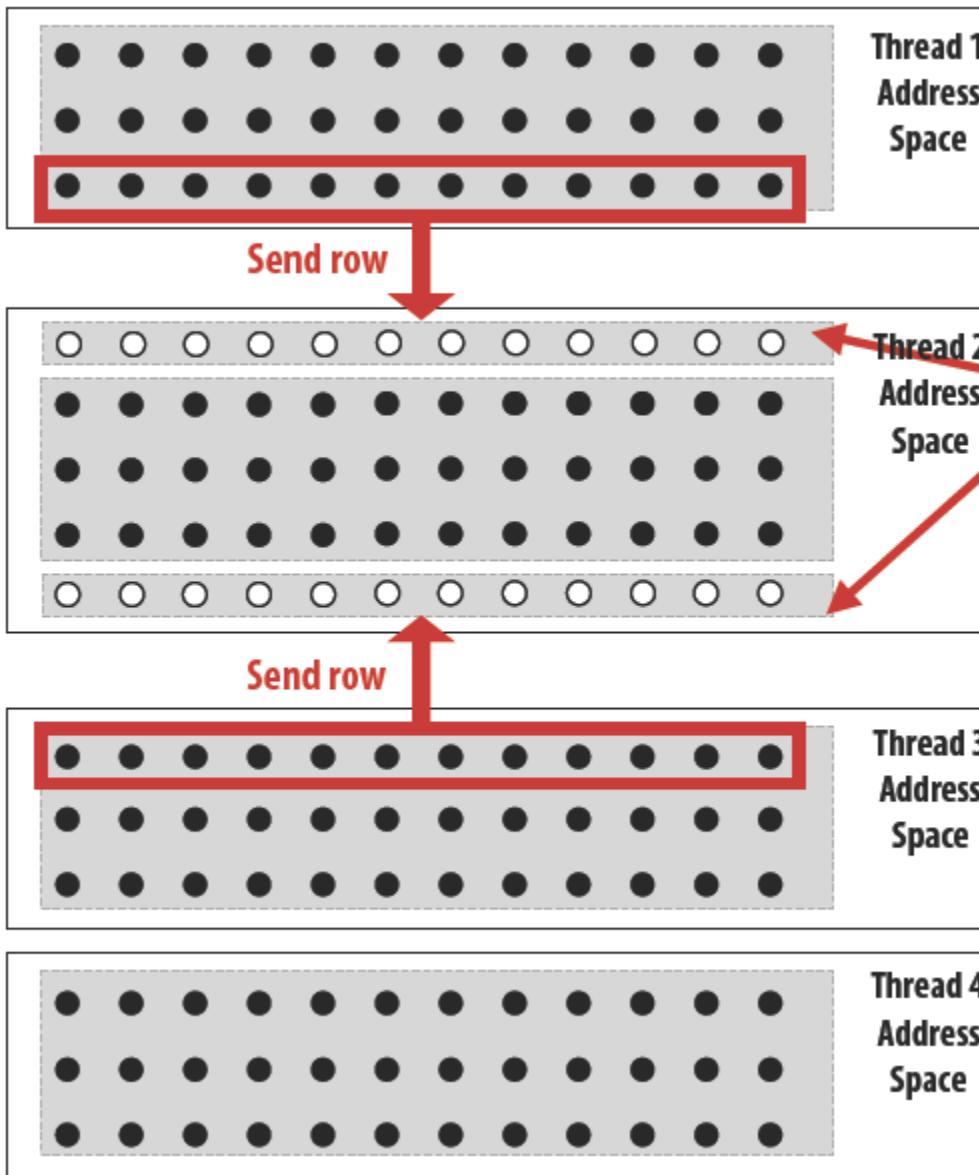


In this figure: four threads

The grid data is partitioned into four allocations, each residing in one of the four unique thread address spaces  
(four per-thread private arrays)

不同的地址空间  
(按线程分配的私有空间)  
(有可能在同一个物理地址空间)

# Data replication is now required to correctly execute the program



## Example:

After red cell processing is complete, thread 1 and thread 3 send row of data to thread 2  
(thread 2 requires up-to-date red cell information to update black cells in the next phase)

*B1 人的数据*  
"Ghost cells" are grid cells replicated from a remote address space. It's common to say that information in ghost cells is "owned" by other threads.

Thread 2 logic: *线程2*

```
float* local_data = allocate(N+2, rows_per_thread+2);  
  
int tid = get_thread_id();  
int bytes = sizeof(float) * (N+2);  
  
// receive ghost row cells (white dots)  
recv(&local_data[0,0], bytes, tid-1);  
recv(&local_data[rows_per_thread+1,0], bytes, tid+1);  
  
// Thread 2 now has data necessary to perform  
// future computation
```

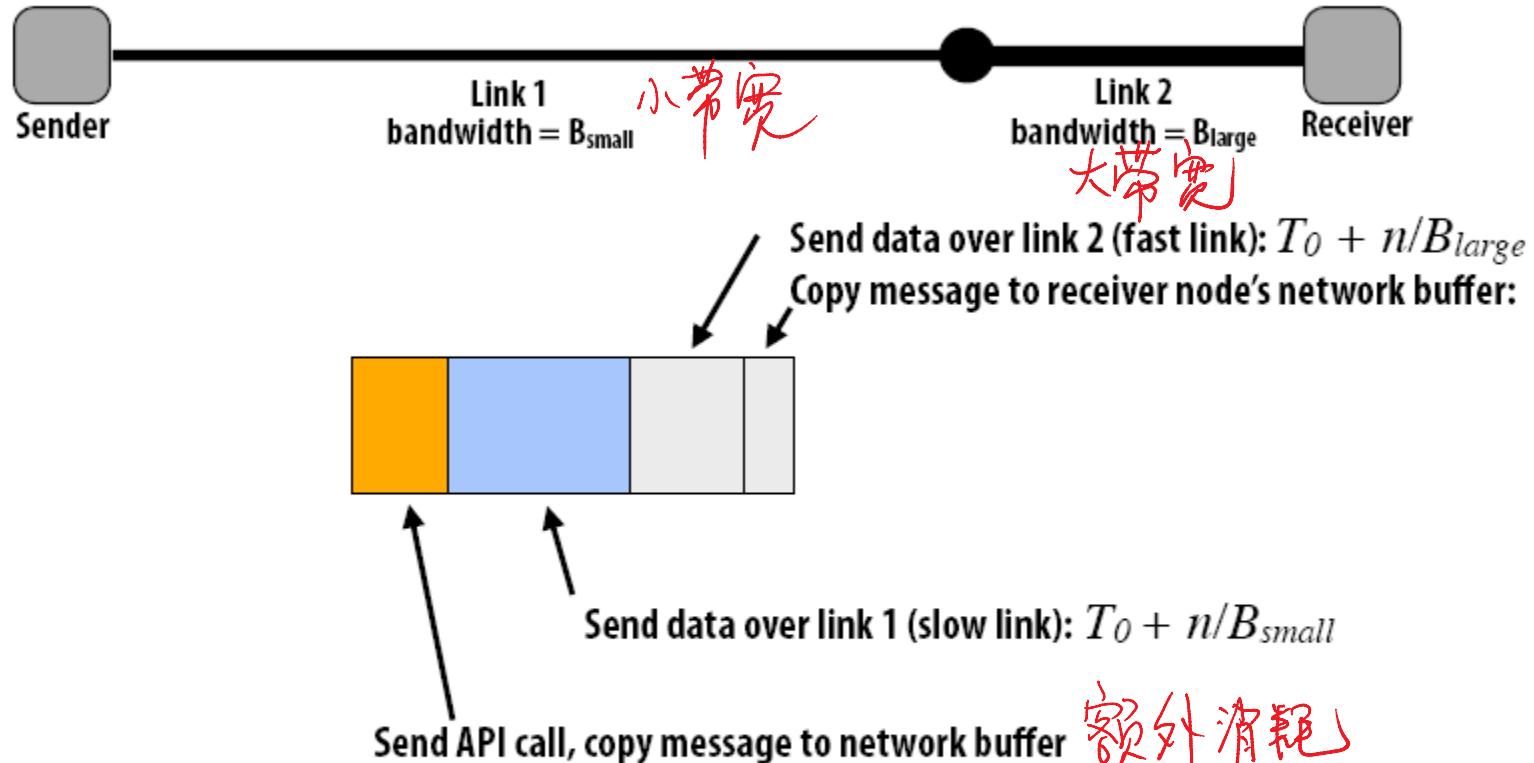
*接收*

# A more general model of communication

Example: sending a n-bit message

更一般的通信模型

Total communication time = overhead + occupancy + network delay



= Overhead (time spent on the communication by a processor)

= Occupancy (time for data to pass through slowest component of system)

= Network delay (everything else)

额外消耗

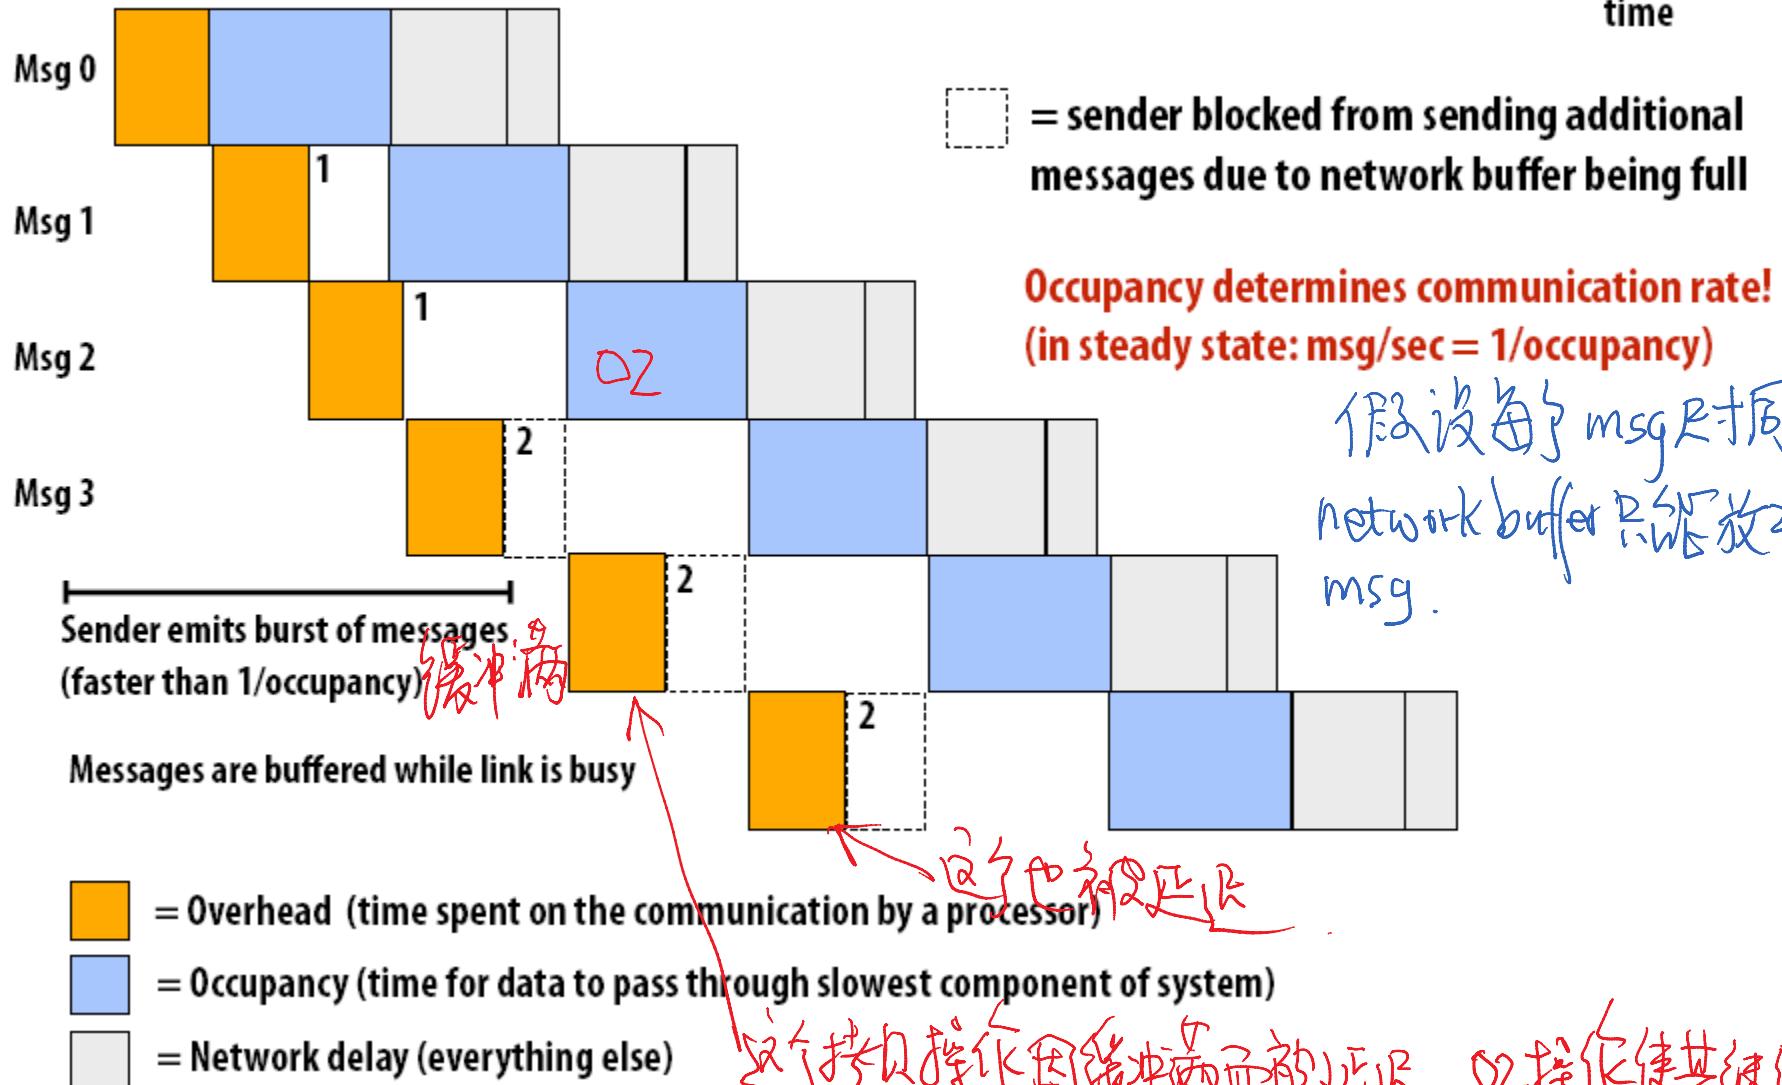
最慢带宽上消耗

其它所有时间

# Pipelined communication

流水式通信

Assume network buffer can hold at most two messages (numbers indicate number of msgs in buffer after insert)



# Think of a parallel system as an extended memory hierarchy

I want you to think of "communication" very generally:

- Communication between a processor and its cache
- Communication between processor and memory (e.g., memory on same machine)
- Communication between processor and a remote memory  
(e.g., memory on another node in the cluster, accessed by sending a network message)

扩展的存储器层次体系

} 更一般地看待  
并行计算系统中的  
通信

View from one processor

Accesses not satisfied in local memory

cause communication with next level

局部存储器无法时,导致下一层通信

So managing locality to reduce the amount of communication performed is important at all levels.

局部性非常重要(全局)

Proc (Core)

Reg

Local L1

Local L2

L2 from another core

L3 cache

Local memory

Remote memory (1 network hop)

Remote memory (N network hops)

低延迟高带宽,小容量

Lower latency, higher bandwidth,  
smaller capacity



高延迟,低带宽,大容量

Higher latency, lower bandwidth,  
larger capacity

# Communication-to-computation ratio

通信计算比

amount of communication (e.g., bytes)

amount of computation (e.g., instructions)

- If denominator is the execution time of computation, ratio gives average bandwidth requirement of code 表示计算所需要的平均带宽
- “Arithmetic intensity” = 1 / communication-to-computation ratio
  - I find arithmetic intensity a more intuitive quantity, since higher is better.
  - It also sounds cooler
- High arithmetic intensity (low communication-to-computation ratio) is required to efficiently utilize modern parallel processors since the ratio of compute capability to available bandwidth is high (recall element-wise vector multiple from previous lecture)

为了有效地利用现代处理器的功能，尽量高的算术强度是需要的！

# Lecture 18. Techniques for reducing communication

- Reduce overhead of communication to sender/receiver
  - Send fewer messages, make messages larger (amortize overhead)
  - Coalesce many small messages into large ones
- Reduce latency of communication
  - Application writer: restructure code to exploit locality
  - Hardware implementor: improve communication architecture
- Reduce contention
  - Replicate contended resources (e.g., local copies, fine-grained locks)
  - Stagger access to contended resources
- Increase communication/computation overlap
  - Application writer: use asynchronous communication (e.g., async messages)
  - HW implementor: pipelining, multi-threading, pre-fetching, out-of-order exec
  - Requires additional concurrency in application (more concurrency than number of execution units)

# Improving temporal locality by fusing loops

合併循環提高時間局部性

```
void add(int n, float* A, float* B, float* C) {  
    for (int i=0; i<n; i++)  
        C[i] = A[i] + B[i];  
}
```

Two loads, one store per math op  
(arithmetic intensity = 1/3)

```
void mul(int n, float* A, float* B, float* C) {  
    for (int i=0; i<n; i++)  
        C[i] = A[i] * B[i];  
}
```

每一個運算有2次讀和1次寫。  
Two loads, one store per math op  
(arithmetic intensity = 1/3)

```
float* A, *B, *C, *D, *E, *tmp1, *tmp2;  
  
// assume arrays are allocated here  
  
// compute E = D + ((A + B) * C)  
add(n, A, B, tmp1);  
mul(n, tmp1, C, tmp2);  
add(n, tmp2, D, E);
```

同上。

Overall arithmetic intensity = 1/3  
總算術強度  $\frac{1}{3}$

```
void fused(int n, float* A, float* B, float* C, float* D, float* E) {  
    for (int i=0; i<n; i++)  
        E[i] = D[i] + (A[i] + B[i]) * C[i];  
}  
  
// compute E = D + (A + B) * C  
fused(n, A, B, C, D, E);
```

Four loads, one store per 3 math ops  
(arithmetic intensity = 3/5)

Code on top is more modular (e.g, array-based math library like numarray/numPy in Python)  
Code on bottom performs much better. Why? 算術強度提高。

# Improve arithmetic intensity by sharing data

共享数据提高浮点强度

- Exploit sharing: co-locate tasks that operate on the same data
  - Schedule threads working on the same data structure at the same time on the same processor
  - Reduces inherent communication
- Example: CUDA thread block
  - Abstraction used to localize related processing in a CUDA program
  - Threads in block often cooperate to perform an operation (leverage fast access to / synchronization via CUDA shared memory)
  - So GPU implementations always schedule threads from the same thread block on the same GPU core

在同一GPU核上调度相同线程内的线程

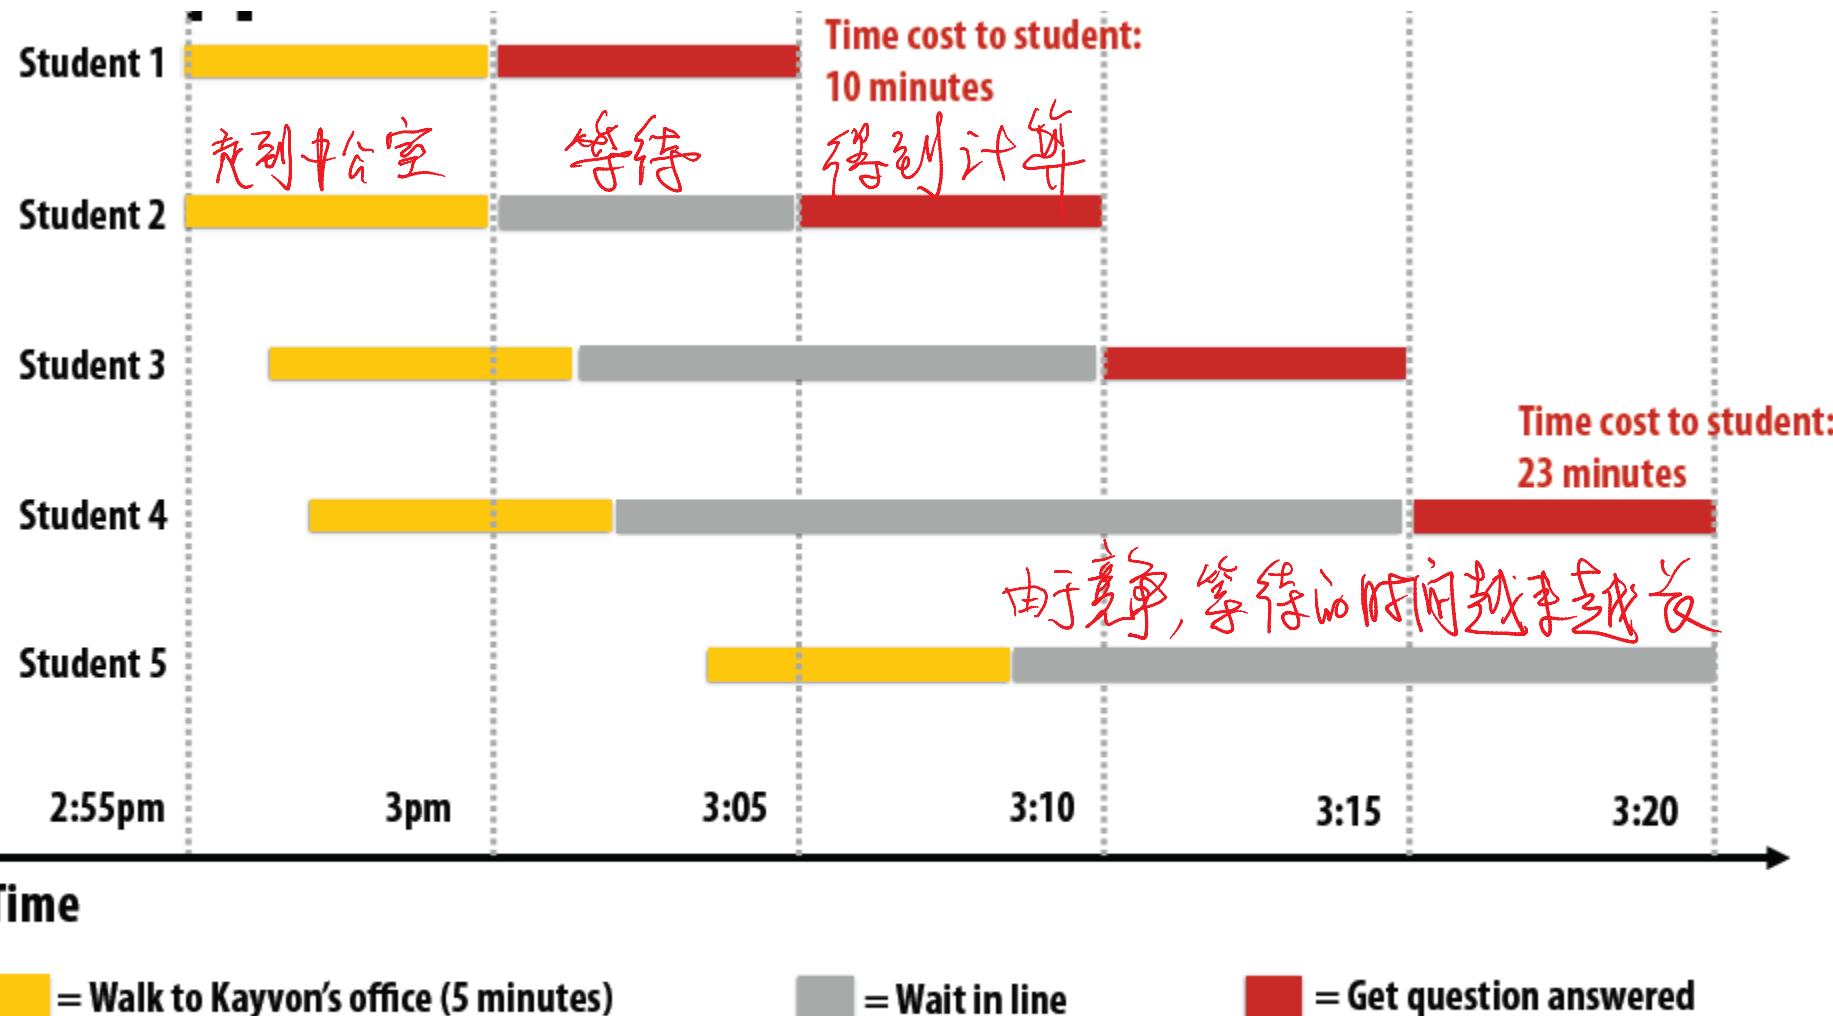
# Exploiting spatial locality

发挥空间局部性.

- **Granularity of communication** can be important because it may introduce artifactual communication
  - **Granularity of communication / data transfer** 通信粒度
  - **Granularity of cache coherence** 数据一致性粒度

缓存一致性粒度

# 共享计算资源



**Problem: contention for shared resource results in longer overall operation times (and likely higher cost to students)**

# Example: Accessing NVIDIA GTX 480 shared memory

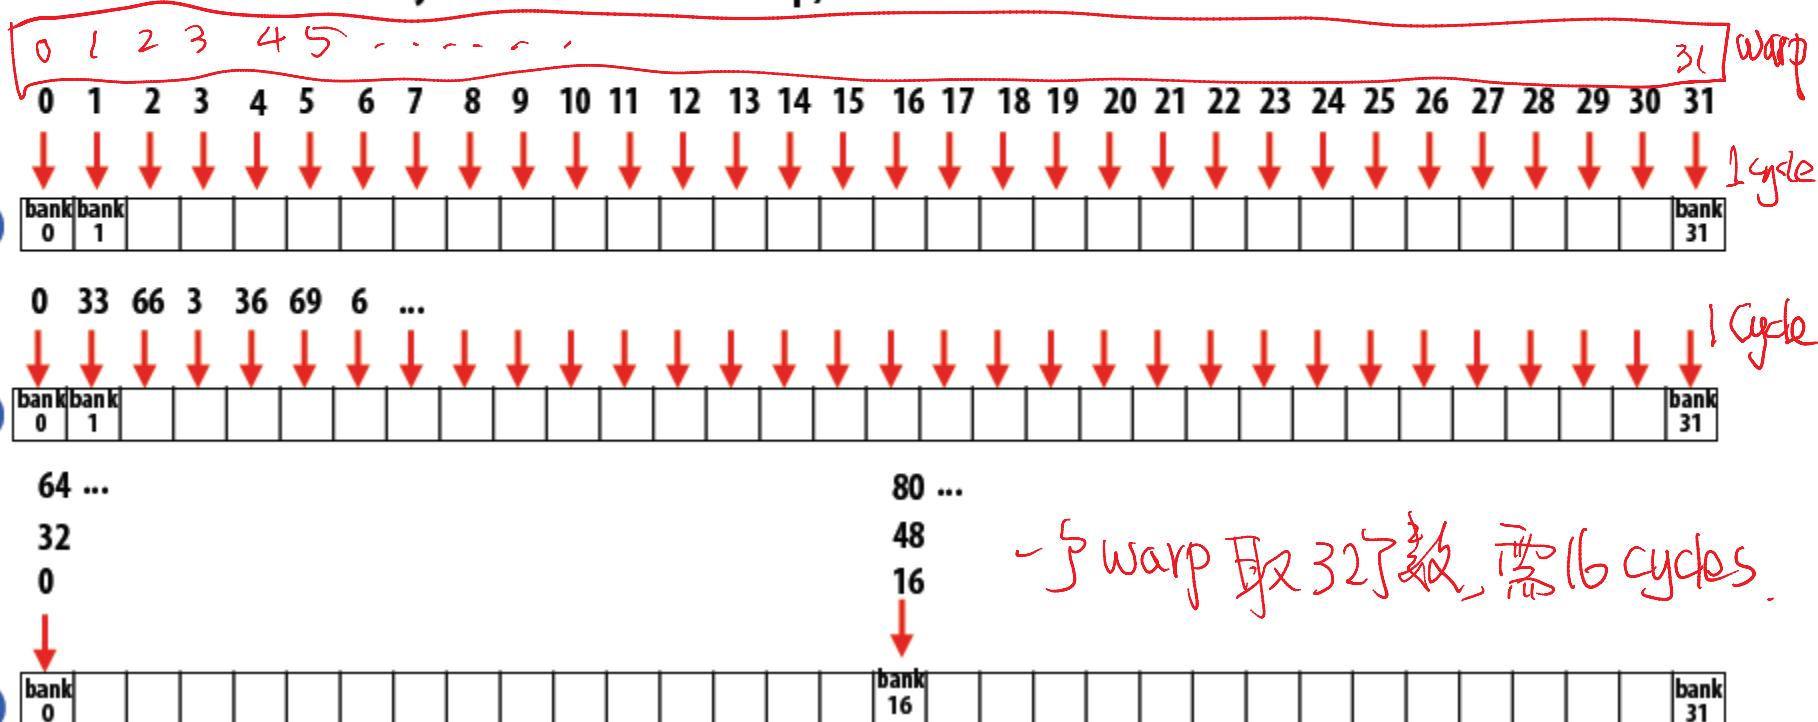
- Shared memory implementation

- On-chip storage, physically partitioned into 32 SRAM banks
- Address X is stored in bank B, where  $B = X \% 32$
- Each bank can provide one word of data to warp per clock

32 threads  
取一个字一个Cycle.

- Figure shows memory addresses requested from each bank as a result of a shared memory load instruction (keep in mind this instruction is executed by all 32 threads in a warp)

```
__shared__ float A[512];  
  
int index = threadIdx.x;  
  
1 float x2 = A[index];      // single cycle  
2 float x3 = A[3*index];    // single cycle  
3 float x4 = A[16 * index]; // 16 cycles
```



# Examination

- Open Book
- Briefly Answer Questions
- Code Parallel programs in a certain programming Model in a certain language with performance Consideration

**Good Luck!**